

Memo Appendix

Figure 1. Average Spread by Fuel Type




Figure 2. Avg Spread by Mileage Band (EVs only)




Figure 3. Avg Spread by Source: High Mileage EVs





Figure 4. Avg Spread by Make & Model: High Mileage EVs



Limitations & Open Questions

This isn't a caveats section to skim past. It's what determines how much weight to put on each finding above, and what I'd want resolved before recommending action on it.

Small sample size

Only 43 rows, and several findings rest on very few cars. The robot versus analyst comparison is based on 6 robot-priced and 8 analyst-priced cases, so the size of the gap is clear but the exact figures are indicative.

Robot vs. analyst case assignment is unverified

Comparing robot against analyst assumes both get similar cases. If analysts are handed the harder ones, part of the gap may come from that rather than accuracy. I could not determine how cases are assigned.

An unexplained flat base price

The script starts every vehicle from the same base price of 305,000 DKK, treating all cars as if they have the same standard value regardless of make, model or fuel type. I did not trace this to the overpricing, but it stood out as worth asking about.

The fix has a lead time

The fix will take time, and until it is in place the engine keeps pricing high mileage EVs the same way. I would want to know whether anything changes in the meantime, or whether we accept the loss while the correction is made.

The monitoring design depends on a genuinely healthy reference period

The control-limit approach only works if the baseline period it learns from is actually healthy. In this dataset, the Tesla overpricing is present from the earliest week available, so a baseline built from it would treat the problem as normal and never raise an alert. This needs to be verified against a known-good period before the monitor goes live.